

Labsheet-6

Kanishka Sharma

CU24250095

Roll No- 26

Section- B

1. Select a suitable dataset such as the House Price Prediction Dataset, Advertising Sales Dataset, or California Housing Dataset.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

housing = fetch_california_housing(as_frame=True)

df = housing.frame

print(df.head())
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	\
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	

	Longitude	MedHouseVal
0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521
3	-122.25	3.413
4	-122.25	3.422

```

MedInc HouseAge AveRooms AveBedrms Population AveOccup Latitude e 8.3252 41.0
6.984127 1.023810 322.0 2.555556 37.88

1      8.3014  21.0 6.238137   o. 971880      2401.o 2.109842 37.86
2      7.2574  52.0 8.288136   1.073446      496.e 2.802260 37.85
3      5.6431  52.0 5.817352 1.073059 558.0 2. 547945 37.85 4 3 .8462 52.0 6.281853 1.081081 565.0
2.181467 37.85

```

```

Longitude MedHouseVa1
o    -122.23      4. 526

1    -122.22  3. 585
2    -122.24  3. 521
3    -122.25  3.413
4    -122.25  3.422

```

2. Import the required libraries including Pandas, NumPy, Matplotlib, and Scikit-learn, and load the dataset.

```

print("Shape of dataset:", df.shape)

print("\nColumns:")
print(df.columns)

print("\nDataset Information:")
print(df.info())

print("\nMissing Values:")
print(df.isnull().sum())

print("\nStatistical Summary:")
print(df.describe())

```

3. Perform basic data preprocessing by handling missing values, encoding categorical variables (if any), and selecting relevant features.

```
print("Missing values before handling:")
print(df.isnull().sum())

df = df.fillna(df.median(numeric_only=True))

print("\nMissing values after handling:")
print(df.isnull().sum())
```

Missing values before handling:

```
MedInc      0
HouseAge    0
AveRooms    0
AveBedrms   0
Population  0
AveOccup    0
Latitude    0
Longitude   0
MedHouseVal 0
dtype: int64
```

Missing values after handling:

```
MedInc      0
HouseAge    0
AveRooms    0
AveBedrms   0
Population  0
AveOccup    0
Latitude    0
Longitude   0
MedHouseVal 0
dtype: int64
```

4. Split the dataset into training and testing sets using an appropriate train-test ratio (e.g., 80:20).

```

X = df.drop("MedHouseVal", axis=1)

y = df["MedHouseVal"]

print("Features:")
print(X.head())

print("\nTarget:")
print(y.head())

```

```

Features:
   MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude  \
0   8.3252     41.0   6.984127   1.023810     322.0   2.555556     37.88
1   8.3014     21.0   6.238137   0.971880    2401.0   2.109842     37.86
2   7.2574     52.0   8.288136   1.073446     496.0   2.802260     37.85
3   5.6431     52.0   5.817352   1.073059     558.0   2.547945     37.85
4   3.8462     52.0   6.281853   1.081081     565.0   2.181467     37.85

   Longitude
0   -122.23
1   -122.22
2   -122.24
3   -122.25
4   -122.25

Target:
0    4.526
1    3.585
2    3.521
3    3.413
4    3.422
Name: MedHouseVal, dtype: float64

```

5. Train a Linear Regression model using the training dataset.

```

X_train, X_test, y_train, y_test =
    train_test_split(x,
                    y,
                    test_size=0.20,
                    random_state=42)

print("Training data: ", X_train.shape)
print("Testing data: ", X_test.shape)
Training data: (16512, 8)
Testing data: (4128, 8)

```

6. Use the trained model to predict values for the testing dataset.

```
model = LinearRegression()

model.fit(X_train, y_train)

print("Model training completed.")
```

Model training completed.

7. Visualize the regression line and compare the actual and predicted values using appropriate plots

```
y_pred = model.predict(X_test)

print("Actual values:")
print(y_test.head())

print("\nPredicted values:")
print(y_pred[:5])
```

Actual values:
20046 0.47700
3024 0.45800
15663 5.00001
20484 2.18600
9814 2.78000
Name: MedHouseVal, dtype: float64

Predicted values:
[0.71912284 1.76401657 2.70965883 2.83892593 2.60465725]

Evaluate the model using performance metrics such as Mean Absolute Error (MAE), Mean Squared Error (MSE), Root Mean Squared Error (RMSE), and R^2 Score.


```
plt.figure(figsize=(8, 6))

plt.scatter(y_test, y_pred, alpha=0.5)

plt.plot(
    [y_test.min(), y_test.max()],
    [y_test.min(), y_test.max()]
)

plt.xlabel("Actual House Values")
plt.ylabel("Predicted House Values")
plt.title("Actual vs Predicted House Values")

plt.show()
```



9. Analyze the regression coefficients to understand the impact of each independent variable on the target variable.

```
mae = mean_absolute_error(y_test, y_pred)

mse = mean_squared_error(y_test, y_pred)

rmse = np.sqrt(mse)

r2 = r2_score(y_test, y_pred)

print("Mean Absolute Error (MAE):", mae)
print("Mean Squared Error (MSE):", mse)
print("Root Mean Squared Error (RMSE):", rmse)
print("R2 Score:", r2)
```

```
Mean Absolute Error (MAE): 0.5332001304956553
Mean Squared Error (MSE): 0.5558915986952444
Root Mean Squared Error (RMSE): 0.7455813830127764
R2 Score: 0.5757877060324508
```

10. Interpret the results and discuss the suitability of Linear Regression for the selected prediction problem.

```
coefficients = pd.DataFrame({
    "Feature": X.columns,
    "Coefficient": model.coef_
})

coefficients["Absolute_Coefficient"] = coefficients["Coefficient"].abs()

coefficients = coefficients.sort_values(
    by="Absolute_Coefficient",
    ascending=False
)

print(coefficients)
```

	Feature	Coefficient	Absolute_Coefficient
3	AveBedrms	0.783145	0.783145
0	MedInc	0.448675	0.448675
7	Longitude	-0.433708	0.433708
6	Latitude	-0.419792	0.419792
2	AveRooms	-0.123323	0.123323
1	HouseAge	0.009724	0.009724
5	AveOccup	-0.003526	0.003526
4	Population	-0.000002	0.000002

	Feature	Coefficient	Absolute_Coefficient
3	AveBedrms	e. 783145	e. 783145
	MedInc	e. 448675	448675
7	Longitude	-0.433708	0.433708
6	Latitude	-0.419792	0.419792
2	AveRooms	-e. 123323	e. 123323
	1 HouseAge	o . 009724	o . 009724
5	AveOccup	-o . 003526	0.003526
			0.000002
4	Population	-0.000002	